

CAPGEMINI 2027 – AI LITERACY MOCK TEST

Engineering Students – Practice Assessment
Section: AI Literacy

Total Questions: 50 MCQs

Suggested Duration: 45–50 Minutes

Difficulty: Moderate to Advanced

Question Type: Scenario-Based Multiple Choice Questions

Options: A, B, C, D

SECTION 1 – AI FOUNDATIONS & GENERATIVE AI

Questions 1–12

Q1. A company deploys a generative AI assistant to summarize incident reports. The assistant produces fluent summaries but occasionally combines details from two different incidents into one account. Which explanation is most appropriate?

- A. Generative AI always merges documents when they contain similar words.
- B. The model may generate a plausible response without reliably preserving the identity and boundaries of each source.
- C. Incident reports cannot be summarized by language models.
- D. Fluency guarantees that the summary is factually consistent.

Q2. A foundation model is adapted for a specialized engineering support application using a smaller collection of domain-specific examples. What is the primary purpose of this adaptation?

- A. To replace all model training with database indexing.
- B. To make the model independent of user instructions.
- C. To improve performance for the target domain or task.
- D. To guarantee that the model will never hallucinate.

Q3. An AI system generates a design explanation that is technically plausible but uses a standard that was withdrawn several years ago. Which control would most directly reduce this risk?

- A. Increase the length of the response.
- B. Ask the model to sound more certain.
- C. Use current authoritative sources and verify time-sensitive claims.
- D. Remove technical terminology from the prompt.

Q4. A language model is asked to answer a question containing an ambiguous abbreviation used differently by two engineering departments. What should the user do to improve the response?

- A. Provide the department context and define the intended abbreviation.
- B. Ask for the longest possible answer.
- C. Repeat the same ambiguous prompt several times.
- D. Assume the model will always infer the correct meaning.

Q5. A student asks an AI assistant to solve a circuit problem, but the assistant assumes an ideal component even though the question specifies non-ideal behavior. What is the best interpretation?

- A. The model has necessarily calculated the problem correctly.
- B. The prompt may have been misunderstood or an assumption may have been introduced incorrectly.
- C. Non-ideal components cannot be represented in language.
- D. The student should accept the assumption because it simplifies the calculation.

Q6. A company uses an LLM to draft compliance responses. The responses are grammatically strong but sometimes omit mandatory exceptions. Which approach is most reliable?

- A. Ask the model to write more confidently.
- B. Use a checklist of required elements and validate each response against approved policy.
- C. Remove exception clauses from the source documents.
- D. Use only the model's general training knowledge.

Q7. A prompt asks an AI system to “make the report better.” The user actually wants fewer words, a formal tone, and stronger evidence. Which change would most improve the prompt?

- A. Add precise criteria describing the desired improvement.
- B. Remove the report from the prompt.
- C. Ask for a creative response without constraints.
- D. Repeat the same phrase many times.

Q8. A developer asks an AI coding assistant to optimize a function. The resulting code is faster on the sample input but fails on large inputs. What should happen next?

- A. Deploy it because the benchmark improved once.
- B. Review complexity, test at scale, and compare behavior against requirements.
- C. Assume the AI has optimized all edge cases.
- D. Remove the large-input tests.

Q9. A user supplies three examples of the required output format before asking the model to process a fourth case. Which technique is being used?

- A. Zero-shot prompting
- B. Few-shot prompting
- C. Model pruning
- D. Data encryption

Q10. A prompt includes a role, objective, constraints, evaluation criteria, examples, and output schema. What is the main advantage?

- A. It guarantees factual accuracy regardless of the task.
- B. It gives the model clearer expectations and reduces ambiguity.
- C. It permanently changes the model's parameters.
- D. It eliminates the need for testing.

Q11. A 200-page technical document is summarized in one pass, but important definitions near the middle are missing. Which workflow is most appropriate?

- A. Divide the document into meaningful sections, summarize them, and verify the combined result.
- B. Ask the model to invent missing definitions.
- C. Remove headings before processing.
- D. Use a shorter prompt with no document context.

Q12. A model receives a prompt containing more information than it can effectively use. Which outcome is possible even if the input is technically accepted?

- A. The model may underuse or lose track of some information.
- B. The model automatically stores every detail permanently.
- C. The model becomes a relational database.
- D. The model must reject every long input.

SECTION 2 – PROMPT ENGINEERING & AI PRODUCTIVITY

Questions 13–24

Q13. A student asks an AI assistant to compare four research papers. The response cites only two papers and appears to mix findings. What is the best first step?

- A. Assume all four papers support the same conclusion.
- B. Map each claim to the correct paper and verify the cited passages.
- C. Ask the model to increase confidence.
- D. Remove the citations to make the answer simpler.

Q14. An AI assistant invents a plausible journal article, complete with author names and a DOI-like string. What does this most clearly demonstrate?

- A. Tokenization
- B. Hallucination
- C. Classification
- D. Fine-tuning

Q15. A policy chatbot receives a question that is not supported by its approved documents. Which behavior is safest?

- A. Provide the most likely answer from general knowledge.
- B. State that the available sources do not support an answer and route the question if necessary.
- C. Create a citation that looks realistic.
- D. Use an unrelated policy as evidence.

Q16. An AI assistant gives two different answers to the same question after small wording changes. What should the user conclude?

- A. At least one answer should be treated as needing verification.
- B. The longer answer is automatically correct.
- C. The latest answer must be correct.
- D. Variation proves that neither answer can ever be useful.

Q17. A company asks an AI system to generate customer refund decisions. Some cases involve legal exceptions and high-value claims. What is the most appropriate design?

- A. Allow fully automatic decisions for all cases.
- B. Use AI for assistance, with policy checks and human review for high-risk or exceptional cases.
- C. Hide the model's uncertainty from reviewers.
- D. Remove exception handling to simplify deployment.

Q18. An employee pastes a confidential architecture diagram and credentials into a public chatbot. Which concern is most serious?

- A. The response may be too short.
- B. Confidential information and credentials may be exposed or mishandled.
- C. The diagram may lose its original colors.
- D. The chatbot may not recognize the file name.

Q19. A company wants to use an AI tool with customer data. Which control is most appropriate before adoption?

- A. Confirm approved data handling, retention, access, and contractual requirements.
- B. Upload a small amount of real customer data without review.
- C. Assume all public tools provide enterprise confidentiality.
- D. Remove only the customer names and ignore other identifiers.

Q20. An AI hiring model shows similar overall accuracy across groups but rejects qualified candidates from one group at a higher rate. What should the organization do?

- A. Treat overall accuracy as sufficient.
- B. Examine group-specific error rates, selection outcomes, and possible causes.
- C. Increase the model's response length.
- D. Remove demographic analysis from evaluation.

Q21. Historical hiring data reflects unequal access to opportunities. A model trained on it reproduces the same pattern. What does this illustrate?

- A. More historical data automatically removes bias.
- B. Training data can encode and reproduce historical inequities.
- C. Bias exists only in prompts.
- D. A model cannot learn patterns from historical records.

Q22. A team evaluates an AI system only on examples it saw during development. Why is this insufficient?

- A. It may not reveal how the system generalizes to unseen situations.
- B. Training examples are always irrelevant.
- C. It guarantees a low score.
- D. It tests privacy more effectively than separate security testing.

Q23. An AI model performs well on average but fails badly on rare safety-critical cases. Which evaluation principle is most relevant?

- A. Evaluate only the average score.
- B. Evaluate important subgroups, edge cases, and high-impact failure modes.
- C. Remove rare cases from the test set.
- D. Prefer fluent outputs over correct outputs.

Q24. A company wants to reduce bias in an image-generation system. Which action is most useful?

- A. Inspect training data and outputs for stereotypes and unequal representation.
- B. Increase image resolution only.
- C. Remove all evaluation criteria.
- D. Ask the model to produce longer captions.

SECTION 3 – CONTEXT, RAG & AI AGENTS

Questions 25–37

Q25. A model's output changes after the underlying data pipeline is modified. What should a responsible team do?

- A. Track the change, test the new pipeline, and compare performance and risks.
- B. Assume the model is unaffected by data changes.
- C. Delete the evaluation records.
- D. Deploy immediately because the change was automated.

Q26. A user wants an AI assistant to return exactly three fields: risk, evidence, and recommended action. Which prompt element is most important?

- A. Output schema or format specification
- B. More decorative language
- C. Model retraining instruction
- D. Screen resolution

Q27. A RAG assistant retrieves documents that contain the correct policy but also several outdated versions. What is the key risk?

- A. The model may combine conflicting versions or select the wrong policy.
- B. Retrieval always improves accuracy regardless of document quality.
- C. The model will automatically delete old documents.
- D. Outdated documents cannot affect generation.

Q28. A company updates its safety procedure but the retrieval index is not refreshed. What is the likely result?

- A. The assistant may continue grounding answers in the older procedure.
- B. The model automatically learns the update from the file name.
- C. The context window becomes unlimited.
- D. The assistant will always refuse to answer.

Q29. A RAG system retrieves relevant passages but the generated answer adds a requirement not present in any passage. Which improvement is most appropriate?

- A. Require claim-level grounding and allow the system to abstain when evidence is insufficient.
- B. Increase creativity settings without evaluation.
- C. Remove source citations.
- D. Ask the model to infer additional requirements.

Q30. A university chatbot must answer questions using only approved course regulations. Which architecture best supports this requirement?

- A. Retrieval from approved documents with source attribution and unsupported-answer handling.
- B. A model that relies only on general pretraining.
- C. A system that retrieves random web pages.
- D. A chatbot that suppresses document metadata.

Q31. An AI agent receives a goal, creates a plan, calls tools, observes results, and revises the next step. What capability is demonstrated?

- A. Static text completion
- B. Goal-directed multi-step action
- C. File compression
- D. Simple keyword matching

Q32. An agent can send emails, update records, and approve payments. Which design principle is most important?

- A. Give it unrestricted access to maximize autonomy.
- B. Use least privilege, approval gates, logging, and monitoring.
- C. Allow it to modify its own permissions.
- D. Disable authentication for internal users.

Q33. A travel-booking agent finds a suitable flight but the fare changes before purchase. What should the agent do?

- A. Complete the booking silently at any price.
- B. Re-check the price and request confirmation when the change affects the user's decision.
- C. Assume the earlier price remains valid.
- D. Cancel all future searches.

Q34. An AI agent is asked to prepare a market report. It searches, compares sources, identifies gaps, and drafts a report. Which step is most important before final delivery?

- A. Verify important claims and ensure the report reflects the requested scope.
- B. Remove all source links.
- C. Assume the first retrieved source is authoritative.
- D. Allow the agent to publish without review.

Q35. A coding assistant generates a database query that runs successfully but exposes records belonging to other users. What failed?

- A. Syntax parsing only
- B. Authorization and data-access reasoning
- C. Tokenization
- D. Text summarization

Q36. A developer asks AI to review a payment-processing function. The assistant reports no vulnerabilities, but no security tests were run. What is the correct conclusion?

- A. No vulnerabilities exist.
- B. The result is only one input to security assurance and requires independent testing.
- C. AI review replaces penetration testing.
- D. Payment code cannot be reviewed by humans.

Q37. A model achieves very high training accuracy and substantially lower validation accuracy. Which explanation is most likely?

- A. Overfitting
- B. Encryption
- C. Retrieval
- D. Tokenization

SECTION 4 – RESPONSIBLE AI, SECURITY & EVALUATION

Questions 38–50

- Q38.** A dataset contains mislabeled examples concentrated in one category. What may happen?
- A. The model may learn misleading patterns and perform poorly for that category.
 - B. The model automatically corrects all labels.
 - C. The model becomes unbiased by default.
 - D. The dataset becomes more representative.
- Q39.** A model is tested on data that was accidentally included in training. Why is the test result misleading?
- A. Data leakage can make performance appear better than it is on truly unseen data.
 - B. Training data always produces lower scores.
 - C. Leakage affects only image models.
 - D. The model cannot use duplicate examples.
- Q40.** A company evaluates an AI assistant before broad deployment. Which evaluation set is most complete?
- A. Fluency only
 - B. Speed only
 - C. Accuracy, robustness, bias, privacy, security, reliability, and human-oversight requirements
 - D. Number of words generated
- Q41.** An attacker places instructions inside a retrieved document telling the AI agent to ignore its system rules and send confidential files. What threat is this closest to?
- A. Prompt injection through external content
 - B. Model compression
 - C. Overfitting
 - D. Data normalization
- Q42.** A company connects an AI agent to internal tools. Which control best limits the impact of a compromised or mistaken agent?
- A. Broad administrator access
 - B. Least-privilege permissions and restricted tool scopes
 - C. No audit logs
 - D. Automatic approval of every action
- Q43.** An AI assistant produces a numerical forecast used in a board presentation. The number differs from the source spreadsheet. What should the employee do?
- A. Use the AI number because it is more recent-looking.
 - B. Reconcile the number against the authoritative data and correct or explain the difference.
 - C. Remove the source spreadsheet.
 - D. Ask AI to make the number more realistic.
- Q44.** A manager uses AI to rank employees using attendance, project results, and unrelated personal information. What is the main concern?
- A. The model may rely on irrelevant or unfair factors.
 - B. More factors always improve fairness.
 - C. Personal information is always relevant to performance.
 - D. AI rankings do not require criteria.
- Q45.** A student uses AI to prepare an explanation that conflicts with official technical documentation. Which source should take priority?
- A. The longer AI response
 - B. The more confident AI response
 - C. The official documentation, subject to checking version and context
 - D. The response with more examples

Q46. An organization wants employees to use AI safely and productively. Which policy is strongest?

- A. Ban all AI tools.
- B. Permit unrestricted use of any public chatbot.
- C. Define approved tools, data rules, verification, training, security controls, and escalation paths.
- D. Allow AI to make all high-impact decisions independently.

Q47. A model is accurate in a laboratory test but performs poorly after deployment because user inputs differ from the test data. What should the team investigate?

- A. Distribution shift and whether evaluation reflects real-world use.
- B. Whether the model should generate longer answers.
- C. Whether all test data should be deleted.
- D. Whether deployment makes models incapable of learning patterns.

Q48. A company wants an AI assistant to answer only when evidence is sufficient. Which behavior should be explicitly designed and evaluated?

- A. Confident guessing
- B. Abstention or escalation when evidence is missing or conflicting
- C. Maximum response length
- D. Automatic acceptance of retrieved text

Q49. An AI-generated report includes a correct conclusion but unsupported intermediate reasoning. Why does this matter?

- A. Unsupported reasoning can make the conclusion difficult to trust or reproduce.
- B. Correct conclusions never need evidence.
- C. Intermediate reasoning is always irrelevant.
- D. A fluent explanation is sufficient proof.

Q50. Before deploying an AI system that affects customers, what is the most appropriate final step?

- A. Conduct documented testing, review risks, define monitoring, and establish human accountability.
- B. Deploy first and investigate only after complaints.
- C. Measure only response speed.
- D. Remove all human escalation routes.

ANSWER KEY

Q Ans. Topic / Concept

1	B	AI reliability and source boundaries	2	C	Foundation models and adaptation
3	C	Knowledge freshness and verification	4	A	Context and ambiguity
5	B	Reasoning limitations	6	B	Structured validation
7	A	Prompt specificity	8	B	AI-assisted coding
9	B	Few-shot prompting	10	B	Structured prompting
11	A	Long-document processing	12	A	Context window limitations
13	B	Source verification	14	B	Hallucination
15	B	Abstention and grounding	16	A	Output variability
17	B	Human oversight	18	B	Data privacy and security
19	A	Data governance	20	B	Fairness evaluation
21	B	Training-data bias	22	A	Evaluation methodology
23	B	Risk-based evaluation	24	A	Bias mitigation
25	A	Model monitoring	26	A	Output formatting
27	A	RAG source quality	28	A	RAG freshness
29	A	Grounded generation	30	A	Grounded AI design
31	B	AI agents	32	B	Agent security
33	B	Agent reliability and confirmation	34	A	Agent workflow validation
35	B	Secure coding	36	B	Security assurance
37	A	Machine learning fundamentals	38	A	Data quality
39	A	Data leakage	40	C	Comprehensive AI evaluation
41	A	Prompt injection	42	B	Least privilege
43	B	Numerical verification	44	A	Relevance and fairness
45	C	Source authority	46	C	Responsible AI governance
47	A	Real-world robustness	48	B	Abstention
49	A	Explainability and evidence	50	A	Deployment governance

Note: This mock test is designed for practice. Candidates should independently verify concepts when using the paper for formal preparation.